

UNIVERSITY OF WATERLOO
Software Engineering

Application of Machine Learning Models as a Solution
to Hyperbola Identification in Ground Penetrating
Radar Data

Sensors & Software Inc.
Mississauga, ON

Prepared by
Xiang Yi (Irene) Chen
Student ID: 20707344
User ID: xy29chen
1B Software Engineering
August 17, 2018

Xiang Yi (Irene) Chen
277 Lester Street
Waterloo, ON N2L 3W6

August 17, 2018

Dr. Derek Rayside, Director
Software Engineering
University of Waterloo
Waterloo, ON N2L 3G1

Dear Dr. Derek Rayside:

The following work term report, entitled “Application of Machine Learning Models as a Solution to Hyperbola Identification in Ground Penetrating Radar Data”, has been prepared for Sensors & Software Inc. as my first work term report for my 1B term. The objective of this work term report is to explore the feasibility of different machine learning models to the problem of hyperbola recognition and classification in ground penetrating radar data.

I would like to thank my supervisor, Adam Fazzari for providing support and guidance along the way, as well as my co-workers for furthering my understanding of ground penetrating radars, from usage to data interpretation. Finally I would like to thank Sensors & Software Inc., a major provider of ground penetrating radar systems, for creating this opportunity and setting me up with the dataset and other resources needed to retrain the models.

I hereby confirm that I have received no further help, other than what is mentioned above, in writing this report. I also confirm that this report has not been previously submitted for academic credit at this or any other academic institution.

Sincerely,

Xiang Yi (Irene) Chen
Student ID: 20707344

Executive Summary

The following report investigates the suitability of a small sample of pre-trained machine learning models, in response to the challenge of identifying hyperbolas in ground penetrating radar (GPR). The problem stems from the fact that hyperbola artefacts in GPR data are highly contextual, and present themselves in many different forms. As such, streamlining the interpretation process with software aids for GPR users proves itself to be a challenge, due to this visual variety in the data.

However, the identification of hyperbola may be solved with the use of machine learning training. Hence, an exploration of convolutional neural networks (CNN), followed by a comparison of three notable CNN models—the Inception V3 model, as well as MobileNet V2 model, and VGG 19 model—to determine their suitability given restrictions to this identification problem. The choice of model was then tested and ultimately proven by a proof-of-concept application to hyperbola identification.

Based on limitations in labelled datasets, restricted training time as well as training environment, the decision to use TODO was made based on the tradeoff between model distributability versus accuracy and training load. The final proof-of-concept model was able to achieve TODO, and its resulting classifier output was XXmb in size, trained by YY labelled images.

Recommendations about TODO based on these limitations were made, in hopes to further improve the performance of the model.

Disclaimer

The pre-trained models and their retraining modules are publicly available on GitHub. They've been applied with proprietary GPR data provided by Sensors & Software Inc.

Table of Contents

Executive Summary	iii
Disclaimer	iii
Table of Contents	iv
List of Figures	v
1 Introduction	1
1.1 Problem Statement	1
1.2 Proposal of Solution	1
2 Comparison & Analysis	3
2.1 General Procedure	3
2.1.1 Retraining Neural Networks	3
2.1.2 Convolutional Neural Networks (CNN)	3
2.2 Inception V3	5
2.3 MobileNet V2	6
2.4 VGG 19	6
3 Conclusions	8
4 Recommendations	9

List of Figures

Figure 2-1:	A neuron in a convolutional network condensing image inputs into a 3D entity. ref:	4
Figure 2-2:	Example overfitted model classifying between two classes, a case to be avoided. ref:	4
Figure 2-3:	Inception V3's CNN model architecture as described. ref:	5
Figure 2-4:	Factorizing and substituting convolution layers of smaller dimensions to cover the input matrix. ref:	6

1 Introduction

Ground Penetrating Radar (GPR) solutions serve a wide range of uses, from underground utility locating, to pavement or ice thickness mapping, to forensics as well as archaeology. However, because of this varied usage—combined with pre-processing and post-processing—GPR data presents itself in many various ways. Hence, interpretation of GPR data is one of the most error-prone aspects of using GPR solutions.

1.1 Problem Statement

One of the biggest challenges of GPR data interpretation is to avoid false positives, such as extreme ringing caused by metal, as well as interference caused by rebounding air waves off the data collection site.

Sensors & Software Inc.’s software is able to apply velocity fitting and other transformations to GPR data to improve and simplify data interpretation, such as identifying false-positive air waves in GPR data through a calculation of its velocity. However, the first and most limiting step to this process is to allow the software to locate hyperbola—evidence of a disturbance in the soil caused by an object.

1.2 Proposal of Solution

Since machine learning has been making advances in the field of computer vision, it seems to be a fit tool to solve this bottleneck in this interpretation process.

Machine learning, in the broadest of terms, is a subset of computer science which allows programs to improve and "learn" without explicitly being told what exactly to do—the construction of algorithms to make predictions based on recognition of patterns on a data set and to make changes to itself so that its future performance is improved.

Specifically, object detection and localization problems generally involve training feature extraction: associating certain characteristics of the image with different weights of importance based on their reoccurrence. In this problem, a hyperbola feature might consist of its curved nature. In this case, the presence—or lack thereof—of hyperbolas in GPR data will be explored with machine learning solutions, as well as their differentiation from other forms of features in GPR data, such as metal ringing and boundaries in the soil.

The follow section will be analysing the steps in a small sample of machine learning models, as well as the transferability and possibility of integration into existing GPR systems.

The comparison process will be taking the following criteria into account:

- Correctness of learning, based on the core transformations of the architecture of the model
- Ease of being transferred for real-time detection, based on the model's cost of memory and processing on ARM systems

For the sake of quickly creating a proof-of-concept model, transfer learning will be applied to an existing open-source pre-trained model. Transfer learning entails modifying layers of existing mature models to customize the model while retaining its generic base layers. Specifically, the following pre-trained models will be explored, chosen for their current popularity and abundant availability in terms of support and resources:

- Inception V3
- MobileNet V2
- VGG 19

2 Comparison & Analysis

For the comparison of potential machine learning models, the general idea is to apply supervised learning, giving labelled datasets to distinguish hyperbolas in the image, then test the model with sets of unlabelled data of various degrees of resemblance to the training dataset.

2.1 General Procedure

2.1.1 Retraining Neural Networks

In the most intuitive sense, neural networks is an imitation of how the human brain works. Input features are weighted positively or negatively, forming and reinforcing connections between common features similar to neurons forming connections between each other.

Each layer performs a transformation on a single vector input, and the learning process is achieved by chaining these layers together, either cyclically or linearly. Deep learning refers to models with a great number of layers of transformations. The penultimate layer—the bottleneck—summarizes the extracted features, to allow the last layer to perform the classification task.

Finally, weights of a model is a measure of the strength of the reinforcements between connections—equivalent to the output parameters to be learned throughout training.

2.1.2 Convolutional Neural Networks (CNN)

Since for this hyperbola identification problem, GPR data can be easily converted to images, thus CNN will be used for their suitability with image inputs. A CNN is a subset of multi-layer neural networks, characterized by their convolution layer process—the union of integrals, or how much two functions overlap as one passes over the other—usually used for image classification. This overlap is the main idea behind how common features between images is reinforced and extracted. Processing is achieved by obtaining the union of the product of input images as matrices.

The common components of CNN model architecture are as follows: the convolutional layer, ReLU layer, pooling layer, and fully-connected layer.

Pooling is a transformation specific to convolutional neural networks: it combines clusters into a single entity to be used in the next layer, to condense and downsample previous




Figure 2-1: A neuron in a convolutional network condensing image inputs into a 3D entity. **ref:**

operations with the intent to reduce parameters. This step assumes an image input and condenses the input into a 3-dimensional array, a convolutional neuron as seen in Figure 2-1. Max pooling, one of the most common types of pooling, is a downsampling function which takes the largest value from the prior cluster, effectively reducing the volume of data to process. Downsampling is not only used to reduce processing load, but by discarding select information, it prevents the model to be overfitted to the training dataset as in Figure 2-2, hence less accurately solving the real-life problem.




Figure 2-2: Example overfitted model classifying between two classes, a case to be avoided. **ref:**

The convolution process is usually paired with a concatenation step as well as an activation ReLU layer. ReLU stands for rectified linear unit, a process to introduce non-linearity in the CNN, effectively performing $f(x) = \max(0, x)$. This is a transformation on convolution outputs that doesn't affect the size of the matrix.

Finally, the fully connected (FC) layer flattens the transformed image matrix. During this FC layer, the softmax function normalizes the output to be between 0 and 1, as a sigmoid function, similar to a categorical probability distribution. In this case, it should be executed in the end to bring the model closer to indicating the probability and likelihood of

hyperbolas.

2.2 Inception V3

Inception V3, also previously known as GoogleNet, is a CNN of five convolutional layers, along with the typical max pooling and softmax functions. However, it differs from the standard design of CNN models by its successive stacking scattered throughout the model, between the convolutional layers, and is hence recognized as a deep training network.

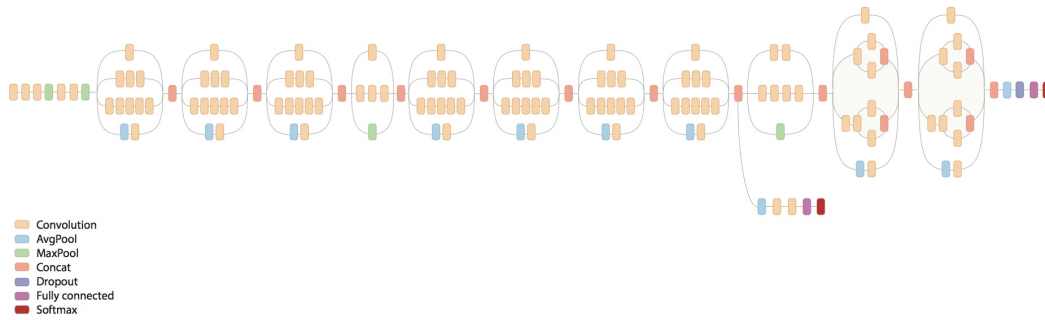


Figure 2-3: Inception V3's CNN model architecture as described. **ref:**

Moreover, Inception extracts features at multiple levels, computing 1×1 , 3×3 , as well as 5×5 convolutional layers, which are concatenated afterwards as seen in Figure 2-3. This particularity of the model allows the training to not only cover greater depths, but allows better receptive field variability of inputs, which in turn leads to being able to detect hyperbolas of different scales when applied. As well, the usage of smaller convolution layers stacked to replace a larger input through factorization, as demonstrated in Figure 2-4, also reduces the process load to train the model. This design is due to the fact that convolutions with larger spacial filters are disproportionately computationally-heavy.




Figure 2-4: Factorizing and substituting convolution layers of smaller dimensions to cover the input matrix. **ref:**

Hence, Inception’s multi-layer convolution process is extremely desirable. Moreover, unlike many CNN models, Inception skips intermediary activation layers, making its computational cost much lower than the other two models compared. This makes Inception a likely candidate to train our hyperbola classifier problem in an environment where memory and computational capacity of the hardware is limited.

As for the result of the model, the weights for Inception V3 are small, coming in at 92MB, with a total depth of 159 stacked layers per training iteration.

2.3 MobileNet V2

MobileNet, as its name suggests, is a CNN model specifically tailored to run on resource constrained environment such as mobile platforms.

To lighten the processing, it applies a depth-wise separable process of convolutions similar to Inception. Essentially, a 3×3 factorized convolutional layer is applied to the input matrix linearly in the depth direction, then a 1×1 convolution concatenates the matrix back layer-wise called a point-wise convolution layer. A total of 14 batches of convolution processes throughout the model before one average pooling function followed by one fully connected layer is applied, ending with the typical softmax function.

However, there is a trade-off in accuracy for this lightweight model.

2.4 VGG 19

VGG 19 consists of five convolutional layers as well, 3×3 stacked on top of each other in increasing depth, designed by the Visual Geometry Group of Oxford. "19" stand for the

number of weight layers in the network, the sum of five batches of convolutional layers as well as max pooling to reduce volume size and finally ending with two fully connected layers as well as a softmax process.

VGG's 3×3 convolutional layers amount to fixed reception fields. Hence, the model is limited to an optimal size of hyperbola to be detected, resulting to over-fitting the model to a certain type of expected GPR data input.

Based on the older AlexNet model—which is characterized by implementing a dropout function after each FC layer—the VGG model attempts to compensate for overfitting with a dropout process after each fully connected layer.

Hence,

advantageous of the VGG architecture over the other two models examined.

However, the downside of this phenomenal accuracy is in its difficult deployment and training process due to the great calculation requirements—a disadvantage both in processing speed and memory requirements.

The network architecture weights themselves are quite large,

Due to its depth and number of fully-connected nodes, the VGG model itself is over 574MB for VGG19. Hence, the biggest drawbacks of VGG, despite its good depth and accuracy, would be hard to deploy onto Sensors and Software Inc's GPR systems, and would only be something usable in EKKO Project, a post-processing software product, if it were running with a powerful GPU.

3 Conclusions

Experimental results show that MobileNet is about 3 times as fast as Inception and 10 times as fast as VGG.

4 Recommendations

My proof-of-concept consists of three training iterations

However, this is but a limited comparison of a small selection of models—though three very successful models.

As well, needless to say, hyperbola detection and identification has proven itself to be a problem difficult to tailor to

at this stage, the current proof-of-concept was trained within a manually created dataset, which is only a rough representation of hyperbola identification. To attempt to capture the full complexity of the problem, a better breadth of images—a wide variety of hyperbolas in different environments, different soil calibrations, different collection speeds and GPR stacking—is required to avoid the model to over-fit to noisy or coincidental features in the GPR data.

No thorough understanding of the complexity of these machine learning models, which is a challenge to adapt and make changes to the CNN. If scaled or applied to the problem naively, certain computational gains during training can be lost, which immediately voids its advantage over other models.